



Level: Beginner

**AWS Certified Cloud Practitioner**[← Back to the Course](#)**Practice Test III – Practice Mode**

Completed on Thu, 26 Oct 2023

2nd  
Attempt

0/65

Marks Obtained



0.00%

Your Score

**FAIL**

Result

**Domain wise Quiz Performance Report**

No.	Domain	Total Question	Correct	Incorrect	Unattempted
1	Cloud Concepts	19	0	0	19
2	Security and Compliance	8	0	0	8
3	Cloud Technology and Services	30	0	0	30
4	Billing, Pricing, and Support	8	0	0	8
<b>Total</b>	<b>All Domains</b>	<b>65</b>	<b>0</b>	<b>0</b>	<b>65</b>

**Review the Answers**

Filter By All Questions

**Question 1**

Unattempted

Domain: Security and Compliance

Which of the following security services can be used to detect users' personal credit card numbers from data stored in Amazon S3?

A. Amazon Macie      right

B. Amazon GuardDuty

C. Amazon Inspector

D. AWS Shield

### Explanation:

#### Correct Answer – A

Explanation: Amazon Macie is a managed security service which can be used to detect personally identifiable information (PII) such as names, password, Credit card numbers from large amounts of data stored in Amazon S3 bucket.

**Option B is incorrect** as Amazon GuardDuty is used to identify threats by analyzing events from AWS CloudTrail, VPC Flow Logs, and DNS Logs. It cannot be used to detect PII from data stored in the Amazon S3 bucket.

**Option C is incorrect** as Amazon Inspector can analyze potential security threats for an Amazon EC2 instance against an assessment template with predefined rules.

**Option D is incorrect** as AWS Shield provides protection against DDOS attacks.

For more information on Amazon Macie, refer to the following URLs:

<https://aws.amazon.com/macie/features/>

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?



### Question 2

Unattempted

**Domain:** Cloud Concepts

Users in the Developer Team need to deploy a multi-tier web application. Which service can be used to create a customized portfolio that will help users for quick deployment?

A. AWS Config

B. AWS Code Deploy

C. AWS Service Catalog right

D. AWS CloudFormation

### Explanation:

**Correct Answer: C**

AWS Service Catalog can be used to create & deploy portfolio of products within AWS infrastructure. This helps to create consistent resources within AWS infrastructure with quick deployment. These catalogues can be used for deployment of single resource or a multi-tier web application consisting of web, application, & database layer resources.

**Option A is incorrect** as AWS config is used for evaluating configuration on the resources deployed in AWS cloud. It will not help for creating portfolios of resources for quick deployment.

**Option B is incorrect** as AWS CodeDeploy is a managed service for automating software deployment on AWS resources & on-premise systems. It is not suitable for creating portfolios of resources for quick deployment.

**Option D is incorrect** as AWS CloudFormation is a service for provisioning AWS resources using templates.

For more information on AWS Service Catalog, refer to the following URL,

<https://aws.amazon.com/servicecatalog/features/>

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?  

**Question 3**

Unattempted

Domain: Cloud Concepts

Which of the following services are used by AWS Service Catalog as a combination to create a portfolio of products?

- A. AWS IAM & AWS Config
- B. AWS Config & AWS CloudFormation
- C. AWS IAM & AWS CloudFormation right
- D. AWS Config & AWS Organizations

**Explanation:****Correct Answer – C**

AWS Service Catalog allows IT organizations to create a portfolio of products that end-users can use to deploy AWS resources as defined in the portfolio. For this, AWS Service Catalog uses AWS IAM & AWS CloudFormation.

**Option A is incorrect** as AWS Service Catalog uses AWS IAM & AWS CloudFormation & not AWS Config to create a portfolio.

**Option B is incorrect** as AWS IAM is used in combination with AWS CloudFormation by AWS Service Catalog.

**Option D is incorrect** as AWS Service Catalog uses neither AWS Config nor AWS Organizations.

For more information on AWS Service Catalog, refer to the following URL;

<https://aws.amazon.com/servicecatalog/features/>

Ask our Experts

 View Queries

Did you like this Question?



#### Question 4

Unattempted

**Domain:** Billing, Pricing, and Support

A client is planning to move some of his On-Premise application workloads to AWS Cloud. He wants to know the elements of AWS cloud for which he will have a cost associated. Which of these elements will NOT require the client to pay for its use?

- A. RDS database backups
- B. Inbound Data transfer      right
- C. EBS Snapshots
- D. Outbound Data Transfer

#### Explanation:

Answer: B

**Option A is incorrect.** RDS provides both automatic backup and user-initiated Point-In-Time snapshots for maintaining reliability. Automatic backups are performed during the backup window of the DB instance as defined during database creation. A storage volume snapshot of the DB instance is created that backs up the entire DB instance and not just the individual databases. RDS database backups are stored in S3 & are priced at \$0.095 per GB month.

**Option B is CORRECT.** AWS works hard for driving down costs allowing customers to reap its benefits. There are no Inbound data transfer charges across all services in all regions. This means that customers can upload petabytes of data to storage destinations like S3 without paying inbound data transfer fees. Data transfers from On-premise to AWS cloud will also not incur Transfer charges.

**Option C is incorrect.** EBS snapshots are stored in S3 that will incur storage costs depending on their size & retention policy of S3. EBS snapshots incur \$0.05 per GB-Month of data stored.

**Option D is incorrect.** Outbound data transfers from AWS services are subject to costs depending on different parameters, like Data transfers from services within AZ's in a region, Data transfers from services across regions, and Data transfers from services to the Internet. Outbound Data transfer prices also vary based on the Region. Data transfers can be costly for Re-hosted applications on AWS that may not be configured or aligned with AWS optimized features. For example, data transfer through a VPC endpoint can be secure & cheaper compared to routing it through the internet. Hybrid cloud architectures need to specifically account for data transfers between the On-Premise & Cloud resources that may become spiky on the billing front.

#### References:

<https://aws.amazon.com/backup/pricing/>

[https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER\\_WorkingWithAutomatedBackups.html](https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_WorkingWithAutomatedBackups.html)

<https://www.cloudmanagementinsider.com/data-transfer-costs-everything-you-need-to-know/>

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>

<https://aws.amazon.com/ebs/pricing/>

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?



#### Question 5

Unattempted

Domain: Cloud Concepts

Which of the following options would entice a company to use AWS over an on-premises data center?  
Choose 2 answers from the options given below.

- A. Having access to Free and Unlimited Storage
- B. Having access to Unlimited Physical servers
- C. Having a highly available infrastructure      right
- D. Ability to use resources on demand      right

#### Explanation:

Answer - C and D

Advantages of the AWS Cloud or any cloud system include the ability to have a highly available infrastructure and the usage of resources on demand.

For more information on the advantages of using AWS, please refer to the below URL:

<https://aws.amazon.com/application-hosting/benefits/>

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?



## Question 6

Unattempted

**Domain:** Security and Compliance

In serverless services such as AWS Lambda, what are the implications of the Shared Responsibility Model?

- A. Amazon has overall responsibility for the infrastructure, including IAM roles and identities that can invoke functions.
- B. The user is responsible for the security and access to the instances that handle the compute capacity.
- C. Amazon is responsible for any malicious code written in the IDE and can terminate any rogue activity.
- D. The user is responsible for IAM roles and identities that can invoke the AWS Lambda functions. right

## Explanation:

### Correct Answer – D

In serverless services such as AWS Lambda, Amazon adopts the responsibility of running all the compute, network and storage of the integrated development environment. The user is responsible for the code itself and identity access management.

<https://aws.amazon.com/lambda/>

**Option A is INCORRECT** because Amazon is not responsible for identity access management (IAM) in the Shared Responsibility Model.

**Option B is INCORRECT** because the user is not responsible for the security and access to the underlying instances that provide compute capacity in serverless services.

**Option C is INCORRECT** because in serverless services such as AWS Lambda, Amazon is not responsible for code written in the IDE.

Ask our Experts

 View Queries

Did you like this Question?



## Question 7

Unattempted

**Domain:** Billing, Pricing, and Support

A large Oil & gas company is planning to deploy a high-volume application on multiple Amazon EC2 instances. Which of the following can help to reduce operational expenses?

- A. Deploy Amazon EC2 instance with Auto-scaling right
- B. Deploy Amazon EC2 instance in multiple AZ's
- C. Deploy Amazon EC2 instance with Amazon instance store-backed AMI
- D. Deploy Amazon EC2 instance with Cluster placement group

## Explanation:

**Correct Answer: A**

Using Amazon EC2 Auto-Scaling helps to match the workload on the application with the optimum number of the Amazon EC2 instance. Due to this, during low load on application, Amazon EC2 instances are terminated which reduces operational cost.

**Option B is incorrect** as deploying an Amazon EC2 instance in a multiple AZ might enhance application availability but will not reduce operational expenses.

**Option C is incorrect** as deploying an Amazon EC2 instance with Amazon instance store-backed AMI incur charges for Amazon EC2 instance usage & storing AMI in Amazon S3. There will be no impact on operational expense using this AMI type.

**Option D is incorrect** as deploying an Amazon EC2 instance in a cluster placement group will help to have low latency between instances but will not reduce operational expenses.

For more information on reducing cost using AWS cloud , refer to the following URL,

<https://aws.amazon.com/economics/>

Ask our Experts

[View Queries](#)

Did you like this Question?



## Question 8

Unattempted

**Domain:** Cloud Concepts

Which of the following are best practices when designing cloud-based systems? (Select TWO)

- A. Build Tightly-coupled components.
- B. Build loosely-coupled components. right
- C. Assume everything will fail. right
- D. Use as many services as possible.

### Explanation:

**Answer – B and C**

Always build components that are loosely coupled. This is so that even if one component does fail, the entire system does not fail.

If you build with the assumption that everything will fail, you will ensure that the right measures are taken to build a highly available and fault-tolerant system.

**Option D is incorrect** because using multiple services increases cost and operational burden, rather use less and efficient services like serverless storage services and serverless compute services.

For more information on a well-architected framework, please refer to the below URL:

<https://docs.aws.amazon.com/wellarchitected/latest/framework/welcome.html>

[Ask our Experts](#)[View Queries](#)

Did you like this Question?



## Question 9

Unattempted

**Domain:** Cloud Concepts

Within the AWS Cloud Adoption Framework (CAF), which of the following correctly represents the four perspectives used to structure the framework's guidance and their associated categories?

- A. Technology, Security, Scalability, Compliance

B. Business, People, Platform, Security right

C. Architecture, Deployment, Performance, Governance

D. Strategy, Implementation, Monitoring, Optimization

## Explanation:

### Correct Answer: B

The AWS Cloud Adoption Framework (CAF) is structured around six perspectives: Business, People, Governance, Platform, Security, and Operations. Each perspective encompasses categories that provide comprehensive guidance for successful cloud adoption.

**Option A is incorrect** because though it contains some relevant concepts, it is not the correct representation of the perspectives and categories in AWS CAF.

**Options C and D are incorrect** because they do not accurately reflect the structure of the AWS CAF.

### Reference:

<https://aws.amazon.com/cloud-adoption-framework/>

Ask our Experts

 View Queries

Did you like this Question?



## Question 10

Unattempted

Domain: Cloud Technology and Services

Which of the following statements related to AWS AppSync is incorrect? (SELECT TWO)

A. AWS AppSync is used to aggregate data from multiple REST APIs / multiple AWS backend data sources(DynamoDB, Aurora, OpenSearch, etc.) in a single API call

B. AWS AppSync is a managed GraphQL service

C. AWS AppSync does not support a strongly typed schema right

D. With AWS AppSync, I can fetch only the required data improving the performance of my application

E. APIs built with AppSync need to use third-party libraries like OpenTelemetry for observability right

## Explanation:

### Correct Answers: C and E

**Option A is incorrect.** AppSync has the ability to aggregate data from multiple data sources, and REST APIs and present it to the end user using a single API call. GraphQL-based API services are based on a single endpoint and the Query Language allows users to specify their queries for representing data using the language.

**Option B is incorrect.** AWS AppSync is a GraphQL-based managed service

**Option C is CORRECT.** AWS AppSync uses the GraphQL schema, which is strongly typed. This allows clients to specify the shape of the data they can request.

**Option D is incorrect.** With the GraphQL language, I can narrow down my results to the exact data that I need. This is helpful in restricting the size of data that an end user will receive thus improving performance

**Option E is CORRECT.** AWS AppSync has Out-of-the-box observability support with AWS CloudWatch and AWS X-Ray

### References:

The following links will help to understand AWS AppSync features in-depth

[AWS AppSync \(amazon.com\)](#)

[Serverless GraphQL and Pub/Sub APIs – AWS AppSync – Amazon Web Services](#)

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?



### Question 11

Unattempted

**Domain:** Cloud Technology and Services

You have been hired as an AWS Architect for a company. There is a requirement to host an application using EC2 Instances. The Infrastructure needs to scale on-demand and also be fault-tolerant. Which of the following would you include in the design? (Select TWO)

A. AWS Auto Scaling right

B. Amazon GuardDuty

C. Elastic Load Balancing right

D. Amazon CloudWatch

## Explanation:

### Correct Answers: A and C

The AWS Documentation mentions the following

You can automatically increase the size of your Auto Scaling group when demand goes up and decrease it when demand goes down. As the Auto Scaling group adds and removes EC2 instances, you must ensure that the traffic for your application is distributed across all of your EC2 instances. The Elastic Load Balancing service automatically routes incoming web traffic across such a dynamically changing number of EC2 instances.

**Option B is incorrect** since Amazon GuardDuty is a threat detection service that continuously monitors for malicious activity and unauthorized behaviour.

**Option D is incorrect** since CloudWatch is a monitoring service

For more information on AWS Autoscaling and ELB, please visit the below URL

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/autoscaling-load-balancer.html>

<https://aws.amazon.com/guardduty/>

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?  

## Question 12

Unattempted

Domain: Cloud Technology and Services

A website for an international sport governing body would like to serve its content to viewers from different parts of the world in their vernacular language. Which of the following services provide location-based web personalization using geolocation headers?

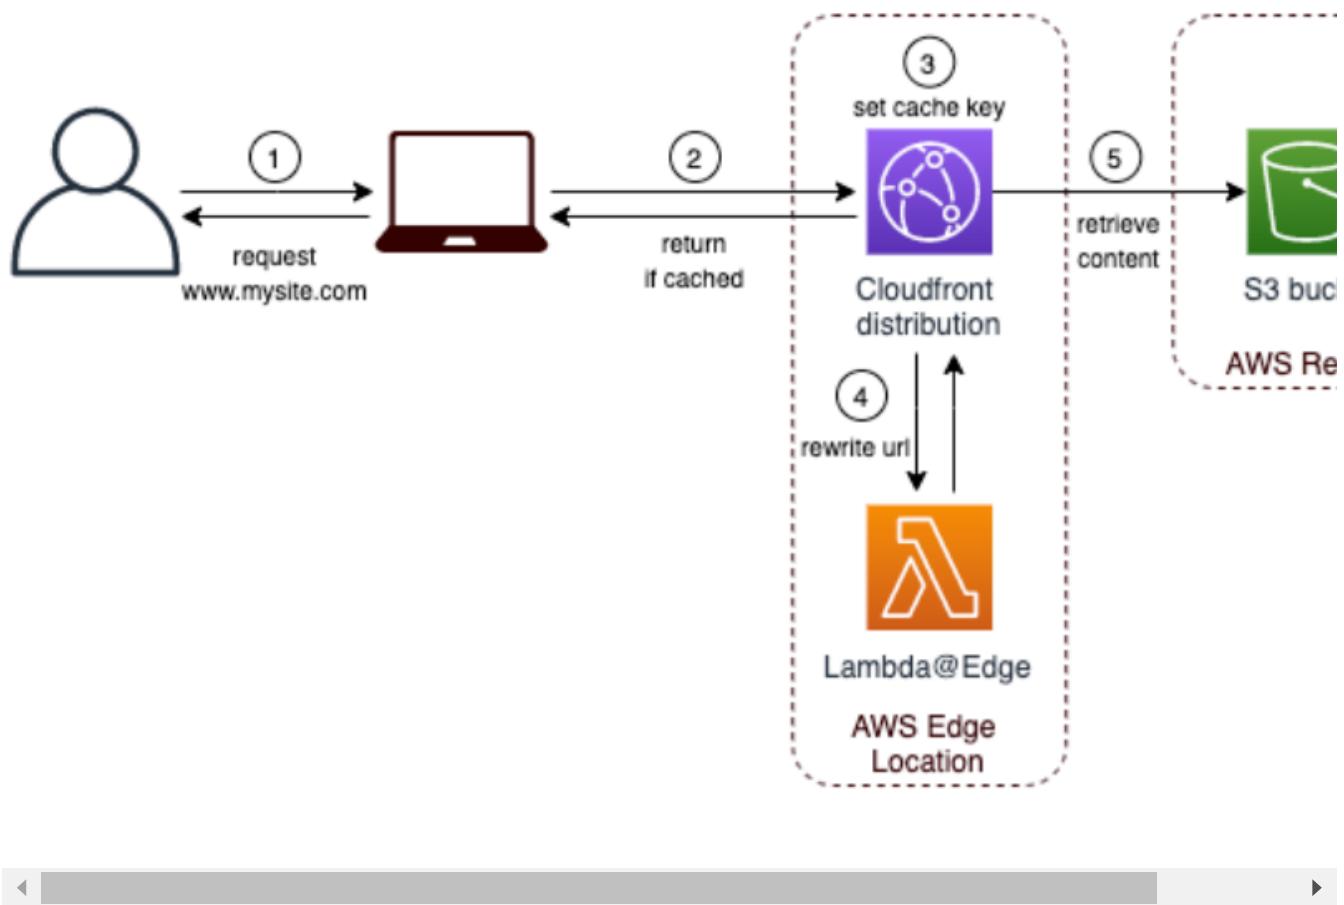
- A. Amazon CloudFront right
- B. Amazon EC2 Instance
- C. Amazon Lightsail
- D. Amazon Route 53

## Explanation:

### Correct Answer – A

Amazon CloudFront supports country-level location-based web content personalization with a feature called Geolocation Headers.

You can configure CloudFront to add additional geolocation headers that provide more granularity in your caching and origin request policies. The new headers give you more granular control of cache behavior and your origin access to the viewer's country name, region, city, postal code, latitude, and longitude, all based on the viewer's IP address.



Option B is INCORRECT because EC2 is just a distractor, not suitable for routing and delivery.

Option C is INCORRECT because Amazon Lightsail will primarily allow for developing, deploying, and hosting websites and web applications. The service will not meet the requirements of the scenario.

Option D is INCORRECT because the geolocation routing policy of Route53 allows different resources to serve content based on the origin of the request. Route 53 does not use geolocation headers.

#### References:

<https://aws.amazon.com/about-aws/whats-new/2020/07/cloudfront-geolocation-headers/>

<https://aws.amazon.com/blogs/networking-and-content-delivery/leverage-amazon-cloudfront-geolocation-headers-for-state-level-geo-targeting/>

[Ask our Experts](#)[View Queries](#)

Did you like this Question?



### Question 13

Unattempted

**Domain:** Cloud Concepts

You have a mission-critical application that must be globally available at all times. If this is the case, which of the below deployment mechanisms would you employ?

- A. Deployment to multiple edge locations
- B. Deployment to multiple Availability Zones
- C. Deployment to multiple Data Centers
- D. Deployment to multiple Regions      right

### Explanation:

Answer – D

Regions represent different geographical locations and are suitable for hosting your application across multiple regions for disaster recovery.

For more information on AWS Regions, please refer to the below URL:

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-regions-availability-zones.html>

[Ask our Experts](#)[View Queries](#)

Did you like this Question?



### Question 14

Unattempted

**Domain:** Security and Compliance

Which of the following can be used to protect against DDoS attacks? Choose 2 answers from the options given below.

- A. AWS EC2
- B. AWS RDS

C. AWS Shield right

D. AWS Shield Advanced right

## Explanation:

Answer – C and D

The AWS Documentation mentions the following:

AWS Shield – All AWS customers benefit from the automatic protections of AWS Shield Standard, at no additional charge. AWS Shield Standard defends against most common, frequently occurring network and transport layer DDoS attacks that target your web site or applications

AWS Shield Advanced – For higher levels of protection against attacks targeting your web applications running on Amazon EC2, Elastic Load Balancing (ELB), CloudFront, and Route 53 resources, you can subscribe to AWS Shield Advanced. AWS Shield Advanced provides expanded DDoS attack protection for these resources.

For more information on AWS Shield, please refer to the below URL:

<https://docs.aws.amazon.com/waf/latest/developerguide/ddos-overview.html>

[Ask our Experts](#)

 View Queries

Did you like this Question?



## Question 15

Unattempted

**Domain:** Cloud Technology and Services

A motion sensor device needs to analyze motion data in its vicinity. Which of the following AWS services can help in the fast processing of data without internet connectivity?

A. IoT Core

B. IoT Analytics

C. IoT Device Management

D. IoT Greengrass right

## Explanation:

**Correct Answer: D**

**Option A is incorrect.** IoT core is categorized as “Control Services” that runs on AWS Cloud. They receive data from IoT devices and send it further for performing analytics e.g. they may send data to Kinesis Firehose and store the data in S3 for further analysis

**Option B is incorrect.** IoT Analytics is categorized as “Data Services” that runs on AWS Cloud & is a fully managed service that makes it easy to run & operationalize sophisticated analytics on massive volumes of IoT data.

**Option C is incorrect.** IoT Device Management is categorized as “Control Services” that runs on AWS Cloud. It is used to track, monitor, and manage device fleets to a large number of devices. It ensures that IoT devices work properly & securely after they have been deployed.

**Option D is CORRECT.** IoT Greengrass provides edge runtime with which IoT applications can be built, deployed & managed on IoT devices. Software built using IoT Greengrass enables devices to act on data locally and can operate even without internet connectivity.

**References:**

The following links will help to understand AWS IoT Greengrass & how AWS IoT works in depth

[What is AWS IoT Greengrass? – AWS IoT Greengrass \(amazon.com\)](#)

[How AWS IoT works – AWS IoT Core \(amazon.com\)](#)

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?

**Question 16**

Unattempted

**Domain:** Cloud Technology and Services

EC2 User Data provides a feature wherein I can do bootstrapping activities when the instance is created. Which of the following statement regarding User Data is incorrect?

- A. User data for an instance can be modified, if the instance is in the stopped state and the root volume is an EBS volume.
- B. User data for an instance can be modified if it is in the running state and the root volume is an EBS volume. right
- C. No need to use 'sudo' in the User Data script to run as the root user.
- D. User Data script can be modified using both the AWS console & the AWS CLI.

**Explanation:**

**Correct Answer: B**

**Option A is incorrect** because it is a true statement. You can modify user data for an instance in the stopped state if the root volume is an EBS volume.

**Option B is correct** because it is a false statement. If the instance is in the stopped state, user data can be modified, not in the running state.

**Option C is incorrect** because it is a true statement. The script commands within User Data executes as root only. So no need to use Sudo explicitly for running them.

**Option D is incorrect** because it is a true statement. User Data script can be modified using AWS CLI also.

**References:**

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/instancedata-add-user-data.html>

[Ask our Experts](#)[View Queries](#)

Did you like this Question?

**Question 17**

Unattempted

**Domain:** Cloud Technology and Services

A web application co-located in two geographically distinct locations is experiencing degraded service in one of the locations. What is the most appropriate routing policy to implement in Amazon Route 53?

- A. Geolocation routing policy
- B. Weighted routing policy
- C. Failover routing policy right
- D. Latency-based routing policy

**Explanation:****Correct Answer – C**

Failover routing policy is the most appropriate routing policy to implement because it will make it possible for traffic to be routed to the resource in good health and not to the one experiencing poor response times. Several instances can be configured.

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html#routing-policy-failover>

**Option A is INCORRECT** because the Geolocation routing policy is essential when traffic to the hosted resources is routed according to its originating source address.

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html#routing-policy-geo>

**Option B is INCORRECT** because a weighted routing policy can distribute request traffic amongst resources. But in this scenario, it would be undesirable and cumbersome to manually set weights (preferences) each time any resource is unhealthy. This will not be the most appropriate routing policy.

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html#routing-policy-weighted>

**Option D is INCORRECT** because latency-based routing policy allows for queries to be served by the resources with the shortest response times to the user's location. The routing policy does not give the required functionality in the scenario.

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html#routing-policy-latency>

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?  

## Question 18

Unattempted

Domain: Cloud Concepts

What is the concept of an AWS region?

- A. It is a collection of Edge locations.
- B. It is a collection of Compute capacity.
- C. It is a geographical area divided into Availability Zones. right
- D. It is the same as an Availability zone.

## Explanation:

Answer – C

A region is a geographical area divided into Availability Zones. Each region contains at least two Availability Zones.

For more information on AWS regions and availability zones, please refer to the below URL-

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-regions-availability-zones.html>

Ask our Experts

 View Queries

Did you like this Question?



### Question 19

Unattempted

**Domain:** Cloud Technology and Services

Users of a popular e-commerce website currently speak to call center agents for query resolution related to the website's products & services. How can the website build virtual contact center agents to enable self-service capabilities for its users quickly?

- A. Using Amazon Polly
- B. Using Amazon Lex right
- C. Using Amazon Personalize
- D. Using Amazon Kendra

### Explanation:

**Correct Answer: B**

**Option A is incorrect.** With Amazon Polly, you can create applications that can talk. It is a service that turns text into speech and builds entirely new categories of speech-enabled products. As an example, your fitness band can tell you that it's now time to re-hydrate.

**Option B is CORRECT.** Amazon Lex is a service for building applications with conversational interfaces using voice and text. A virtual contact center agent resembles a conversational AI for ChatBots answering customer queries thus enabling self-service capabilities for the customer. Amazon Lex helps reduce design time and accelerates Bot deployment from weeks to hours.

**Option C is incorrect.** Amazon Personalize helps provide personalized recommendations for users. It is a fully managed ML service that uses user's data to generate personalized recommendations for the user. As an example, customizable Amazon Personalize resources can be used to add multiple types of personalized product recommendations as seen on e-commerce sites like amazon.com "Recommended for you, frequently bought together"

**Option D is incorrect.** Amazon Kendra is an enterprise search service powered by ML. Search capabilities can be added to applications that will help end-users discover information stored as content within their company

**Reference:**

The following AWS documentation is a blog that describes how Amazon Lex can expedite conversational AI design

[Expedite conversation design with the automated chatbot designer in Amazon Lex | AWS Machine Learning Blog](#)

[Ask our Experts](#)[View Queries](#)

Did you like this Question?

**Question 20**

Unattempted

**Domain:** Security and Compliance

Which of the following can be used to manage identities in AWS?

- A. AWS Config
- B. AWS IAM right
- C. AWS Trusted Advisor
- D. AWS

**Explanation:**

Answer – B

The AWS Documentation mentions the following:

AWS Identity and Access Management (IAM) is a web service that helps you control access to AWS resources securely. You use IAM to control who is authenticated (signed in) and authorized (has permissions) to use resources.

For more information on AWS IAM, please refer to the below URL:

<https://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>[Ask our Experts](#)[View Queries](#)

Did you like this Question?



**Question 21**

Unattempted

**Domain:** Cloud Technology and Services

You are working on an e-commerce web application running on an EC2 instance. But it is experiencing a bad performance in browsing and searching use cases when heavy load use-cases are running simultaneously. The application monitors highlight a bottleneck in the web tier.

You decide to re-engineer the application code by **decoupling** the web tier components from the order's heavy workloads. What AWS service can support the application change?

- A. AWS Auto Scaling
- B. AWS Elastic Load Balancing
- C. Amazon SQS right
- D. Amazon Kinesis Streams

**Explanation:****Answer: C**

**Option A is INCORRECT** because it is not a service for decoupling use-case. AWS Auto Scaling adjusts the capacity to maintain performance leaving unchanged the application design.

**Option B is INCORRECT** because AWS ELB is a service to distribute the workload across EC2 instance, not a service to support application refactoring or re-engineering.

**Option C is CORRECT** because Amazon SQS implements messaging that is a typical integration pattern to decouple application components. AWS documentation mentions "Amazon SQS offers a reliable, highly-scalable hosted queue for storing messages as they travel between applications or microservices. It moves data between distributed application components and helps you decouple these components".

**Option D is INCORRECT** because it is a streaming service not suitable for the scenario.

**Diagram:** none**References:**

<https://aws.amazon.com/sqs/faqs/>

**Ask our Experts****View Queries**

Did you like this Question?



**Question 22**

Unattempted

**Domain:** Cloud Technology and Services

An online streaming company is prohibited from broadcasting its content in certain countries and regions in the world. Which Amazon Route 53 routing policy would be the most suitable in guaranteeing their compliance?

- A. Geoproximity
- B. Geolocation right
- C. Multi-value answer
- D. Failover

**Explanation:****Correct Answer – B**

Amazon Route 53 geolocation routing policy makes it possible for different types of content to be served depending on the browser's geographical location. In this use case, the streaming company can serve a restriction message if Amazon Route 53 detects origin requests from prohibited countries.

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html#routing-policy-geo>

**Option A is INCORRECT** because geo-proximity allows for DNS traffic to be routed in accordance with a bias or preset preference rule. This allows the user to be served with content from resources closest to their geographical location. This routing manipulates DNS traffic flow only. This routing policy is not the most suitable.

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html#routing-policy-geoproximity>

**Option C is INCORRECT** because a multi-value answer primarily addresses the quality of service and resources queried in DNS requests. This routing policy is not the most suitable.

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html#routing-policy-multivalue>

**Option D is INCORRECT** because failover allows for the automatic switch to healthy DNS resources if another becomes unavailable. It will not allow for the preferential serving of content based on the geographical location. This routing policy is not the most suitable.

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html#routing-policy-multivalue>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?



### Question 23

Unattempted

**Domain:** Billing, Pricing, and Support

Which of the following are attributes that determine the cost of On-Demand EC2 Instance? (Select THREE)

- A. Instance Type right
- B. Operating System right
- C. Region right
- D. Edge location

### Explanation:

**Answer- A, B and C**

If you see the below snapshot from the EC2 on-demand pricing page, you can see the different components that make up the pricing.

### On-Demand Plans for Amazon EC2

Select a region, operating system, instance type, and vCPU to view rates

Region

US East (Ohio)

Operating system

Linux with SQL Web

Instance type

All

vCPU

1

For more information on AWS EC2 On-Demand pricing, please refer to the below URL:

<https://aws.amazon.com/ec2/pricing/on-demand/>

[Ask our Experts](#)[View Queries](#)

Did you like this Question?



### Question 24

Unattempted

**Domain:** Cloud Technology and Services

A company wants to utilize AWS storage. For them, low storage cost is paramount. The data is rarely retrieved and a data retrieval time of 13-14 hours is acceptable for them. What is the best storage option to use?

- A. Amazon S3 Glacier
- B. S3 Glacier Deep Archive right
- C. Amazon EBS volumes
- D. AWS CloudFront

### Explanation:

Answer – B

**S3 Glacier Deep Archive** offers the lowest cost storage in the cloud, at prices lower than storing and maintaining data in on-premises magnetic tape libraries or archiving data offsite.

It expands our data archiving offerings, enabling you to select the optimal storage class based on storage and retrieval costs, and retrieval times.

**Option B is correct** because S3 Glacier Deep Archive offers low-cost storage and retrieval time doesn't matter for the company. If the question asks for fast retrieval time then S3 Glacier would be correct.

**Option A is incorrect** because S3 Glacier is not cheaper than S3 Glacier Deep Archive.

**Options C and D are incorrect** because they are not suitable for data archive and faster retrieval. Also, the CloudFront is not for storage.

With **S3 Glacier**, customers can store their data cost-effectively for months, years, or even decades. S3 Glacier enables customers to offload the administrative burdens of operating and scaling storage to AWS, so they don't have to worry about capacity planning, hardware provisioning, data replication, hardware failure detection, and recovery, or time-consuming hardware migrations.

Amazon S3 Glacier for archiving data that might infrequently need to be restored within a few hours

S3 Glacier Deep Archive for archiving long-term backup cycle data that might infrequently need to be restored within 12 hours

Storage class	Expedited	Standard	Bulk
Amazon S3 Glacier	1–5 minutes	3–5 hours	5–12 hours
S3 Glacier Deep Archive	Not available	Within 12 hours	Within 48 hours

Storage class	Expedited	Standard	Bulk

**Reference:**

<https://docs.aws.amazon.com/amazonglacier/latest/dev/introduction.html>

<https://docs.aws.amazon.com/prescriptive-guidance/latest/backup-recovery/amazon-s3-glacier.html>

<https://aws.amazon.com/s3/storage-classes/>

[Ask our Experts](#)

 [View Queries](#)

Did you like this **Question?**

**Question 25**

Unattempted

**Domain:** Cloud Technology and Services

A software development team needs to automate compiling code, packaging it, creating a docker image, and running unit tests for the code they develop. Which AWS service can help them achieve this?

- A. AWS CodeBuild      right
- B. AWS CodeDeploy
- C. AWS CodeCommit
- D. All of the above

**Explanation:**

**Correct Answer: A**

**Option A is CORRECT.** AWS provides a suite of services as DevOps tools for automating the CI/CD process. AWSCodeBuild is a fully managed Continuous Integration service that can compile code, run tests, create Docker images, and produce packages that are ready to deploy

**Option B is incorrect.** AWSCodeDeploy is a managed deployment service that is responsible for automating software deployments to a variety of compute services like EC2, ECS, and Lambda. CodeDeploy takes inputs from packages that CodeBuild produces.

**Option C is incorrect.** AWSCodeCommit is a managed source control service provided by AWS for hosting GIT repositories.

**Option D is incorrect.** CodeBuild will help achieve the necessary requirements provided in the question.

**References:**

The following documentation provides deeper insights into AWSCodeBuild & other DevOps tools provided by AWS

[Managed Build Server - AWS CodeBuild - AWS \(amazon.com\)](#)

[DevOps - Amazon Web Services \(AWS\)](#)

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?

**Question 26**

Unattempted

**Domain:** Cloud Concepts

Which AWS service provides a fully managed NoSQL database service that provides fast and predictable performance with seamless scalability?

- A. AWS RDS
- B. DynamoDB right
- C. Oracle RDS
- D. Elastic Map Reduce

**Explanation:**

Answer: - B

DynamoDB is a fully managed NoSQL offering provided by AWS. It is now available in most regions for users to consume.

For more information on AWS DynamoDB, please refer to the below URL:

<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/introduction.html>

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?



**Question 27**

Unattempted

**Domain:** Billing, Pricing, and Support

In the AWS Billing and Management service, which tool will allow the user to graphically visualize billing and usage over time, particularly monthly running costs?

- A. AWS Bills
- B. AWS Cost Explorer right
- C. AWS Reports
- D. AWS Budgets

**Explanation:****Correct Answer – B**

AWS Cost Explorer allows the user to generate a graphical representation of their billing and usage over time.

<https://aws.amazon.com/aws-cost-management/aws-cost-explorer/>

**Option A is INCORRECT** because AWS Bills will list the historical costs that would have been incurred over the past month with granular options. The tool will not give the graphical visualization as specified in the question.

**Option C is INCORRECT** because AWS Reports will give a composite overview of costs and usage. The tool gives a granular perspective of usage and billing but without a graphical output.

**Option D is INCORRECT** because AWS Budgets will give the user the status of user-set budgets and provide forecasts of estimated costs. The tool will not give a graphical representation of the data.

[Ask our Experts](#)[View Queries](#)

Did you like this Question?

**Question 28**

Unattempted

**Domain:** Cloud Technology and Services

An educational institution has desktop applications that they want to convert to SaaS without rewriting & providing streaming access to students online. Which AWS service will help do so?

- A. Amazon WorkSpaces
- B. Amazon AppStream right

C. Amazon WorkSpaces Web

D. Both B & C

## Explanation:

Correct Answer: B

**Option A is incorrect.** Amazon WorkSpaces is used as a fully managed cloud-based Desktop As A Service (DaaS) environment. So, if you have a software development team that needs Virtual Desktops then, Amazon WorkSpaces can be used to provision them. Amazon WorkSpaces offers the flexibility the software team needs to complete their work while offering a pay-as-you-go model and top-level scalability

**Option B is CORRECT.** Existing Desktop apps can be easily converted to SaaS without rewriting them. This can be done by using Amazon AppStream. These legacy applications can then be streamed and made available to users.

**Option C is incorrect.** WorkSpaces Web is a low cost fully managed Workspace that is built specifically for facilitating web-based workloads (internal websites, SaaS applications). Since the requirement here is only to provide access to SaaS applications, Appstream will be a better choice as compared to WorkSpaces Web

**Option D is incorrect.** Although WorkSpaces Web looks & functions similar to AppStream, with the given requirements, we go with AppStream as the correct answer.

## References:

The following links will help to understand Amazon AppStream features in-depth

[Amazon AppStream 2.0 Features | Non-Persistent Virtualization](#)

[Desktops and Apps with Selective Persistence – Amazon AppStream 2.0 – AWS](#)

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?  

## Question 29

Unattempted

**Domain:** Billing, Pricing, and Support

In the AWS Billing and Management service, which tool can provide usage-based forecasts of estimated billing costs and usage for the coming months?

A. AWS Cost Explorer right

B. AWS Bills

**C. AWS Reports****D. Cost & Usage Reports****Explanation:****Correct Answer – A**

AWS Cost Explorer can create user-defined custom forecasts for future usage patterns.

<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/ce-forecast.html>

<https://aws.amazon.com/about-aws/whats-new/2019/07/usage-based-forecasting-in-aws-cost-explorer/>

**Option B is INCORRECT** because AWS Bills will list the historical costs that would have been incurred over the past month with granular options. The tool will not give the usage-based forecasts as specified in the question.

**Option C is INCORRECT** because AWS Reports will give a composite overview of costs and usage. The tool gives a granular perspective of usage and billing but without usage-based forecasts.

**Option D is INCORRECT** because AWS Reports and Cost & Usage Reports are the same tool. Option C. explanation outlines why it is inaccurate as a response to the question.

[Ask our Experts](#)[View Queries](#)

Did you like this Question?

**Question 30**

Unattempted

**Domain:** Cloud Concepts

Which of the below AWS services allows you to increase the number of resources based on the demand of the application or users?

- A. AWS EC2
- B. AWS Autoscaling right
- C. AWS ELB
- D. AWS Inspector

**Explanation:**

Answer - B

The AWS Documentation mentions the following:

AWS Auto Scaling enables you to configure automatic scaling for the scalable AWS resources for your application in a matter of minutes. AWS Auto Scaling uses the Auto Scaling and Application Auto Scaling services to configure scaling policies for your scalable AWS resources.

For more information on AWS Autoscaling, please refer to the below URL:

<https://docs.aws.amazon.com/autoscaling/plans/userguide/what-is-aws-auto-scaling.html>

[Ask our Experts](#)

 [View Queries](#)

Did you like this **Question?**



### Question 31

Unattempted

**Domain:** Cloud Technology and Services

Which of the following AWS managed database service provides processing power that is up to 5X faster than a traditional MySQL database?

- A. MariaDB
- B. Aurora      right
- C. PostgreSQL
- D. DynamoDB

### Explanation:

Answer – B

The AWS Documentation mentions the following:

Amazon Aurora (Aurora) is a fully managed, MySQL- and PostgreSQL-compatible, relational database engine. It combines the speed and reliability of high-end commercial databases with the simplicity and cost-effectiveness of open-source databases. It delivers up to five times the throughput of MySQL and up to three times the throughput of PostgreSQL without requiring changes to most of your existing applications.

For more information on AWS Aurora, please refer to the below URL:

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Aurora.Overview.html>

[Ask our Experts](#)[View Queries](#)

Did you like this Question?



### Question 32

Unattempted

**Domain:** Cloud Concepts

Which of the following activities is within the scope of AWS Support? (Select TWO)

- A. Troubleshooting API issues      right
- B. Code Development
- C. Debugging custom software
- D. Third-party application configuration on AWS resources      right
- E. Database query tuning

### Explanation:

#### Correct Answers: A and D

As a part of AWS Support following activities are performed,

1. Queries regarding all AWS Services & features.
2. Best Practices to integrate, deploy & manage applications in the AWS cloud.
3. Troubleshooting API & SDK issues.
4. Troubleshooting operational issues.
5. Issues related to any AWS Tools.
6. Problems detected by EC2 health checks
7. Third-Party application configuration on AWS resources & products.

AWS Support does not include:

Code development

Debugging custom software

Performing system administration tasks

Database query tuning

Cross-Account Support

**Option B is incorrect** as Code Development is not in the scope of AWS Support. This needs to be taken care of by the customer.

**Option C is incorrect** as Debugging custom software is not in the scope of AWS Support. This needs to be taken care of by the customer.

**Option E is incorrect** as Database query tuning is not in the scope of AWS Support. This needs to be taken care of by the customer.

For more information on AWS Support, refer to the following URL,

<https://aws.amazon.com/premiumsupport/>

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?



### Question 33

Unattempted

**Domain:** Security and Compliance

Which of the following services helps in governance, compliance, and risk auditing in AWS?

- A. AWS CloudFormation
- B. AWS CloudTrail      right
- C. AWS CloudWatch
- D. AWS SNS

### Explanation:

Answer - B

The AWS Documentation mentions the following:

AWS CloudTrail is a service that enables governance, compliance, operational auditing, and risk auditing of your AWS account. With CloudTrail, you can log, continuously monitor, and retain account activity related to actions across your AWS infrastructure. CloudTrail provides event history of your AWS account activity, including actions taken through the AWS Management Console, AWS SDKs, command line tools, and other AWS services. This event history simplifies security analysis, resource change tracking, and troubleshooting.

For more information on AWS CloudTrail, please refer to the below URL:

<https://aws.amazon.com/cloudtrail/>[Ask our Experts](#)[View Queries](#)

Did you like this Question?



### Question 34

Unattempted

**Domain:** Cloud Technology and Services

My organization has certain software usage policies that I need to comply with. How can I store customized software libraries that my team can use for software development on AWS Cloud?

- A. Using AWS CodeCommit
- B. Using AWS CodeArtifact      right
- C. Using AWS CodeGuru
- D. I cannot customize and store software libraries in AWS. I need to import them directly from public repositories like Maven Central

### Explanation:

**Correct Answer:** B

**Option A is incorrect.** AWS CodeCommit is a managed source control service repository provided by AWS where developers can write code & commit it to these repositories. Code written & committed to CodeCommit will internally use these software libraries that are either downloaded from a public repository like 'Maven Central' or from an Artifact repository like CodeArtifact.

**Option B is CORRECT.** AWS CodeArtifact is a fully managed artifact repository service. It makes it easy for Organizations to securely store, customize, publish & share software libraries that are used by developers during software development. As an example, my Organization may mandate customization of any software library that is present in public repositories like 'Maven Central' prior to its use. Post customization, I would repackage the software library and store it in CodeArtifact for developers to use.

**Option C is incorrect.** AWS CodeGuru is a code vulnerability assessment tool that checks and detects security vulnerabilities within code, and deviations from best practices while using AWS APIs & SDKs and provides recommendations for the same.

**Option D is incorrect.** I can customize my software libraries based on my organization's software policy usage & store them in AWS CodeArtifactory for developers to use during development.

### Reference:

[Ask our Experts](#)[View Queries](#)

Did you like this Question?



### Question 35

Unattempted

**Domain:** Cloud Concepts

A website designed using On-Premise servers has a consistent increase in traffic during certain occasions like Christmas & Thanksgiving. The first year saw a 20% average increase while the second year saw a 40% average increase. It is observed that every year I need to spend on new Servers that lie idle after usage. If I use AWS Cloud resources instead, the advantages that I will get are \_\_\_\_\_.

- A. I can go Global in minutes
- B. I can increase speed & agility of my website
- C. I do not need to maintain the DNS name
- D. I can save on costs due to the elastic nature of the cloud right

### Explanation:

**Correct Answer: D**

As seen from the scenario, every year I need to budget for new servers due to an increase in traffic. This adds to my cost, maintenance of the servers that may not be required at other times in the year. Due to the elastic nature of the cloud, I can directly avail cost benefits by using a pay-as-you-go model and dynamically scaling - OUT & scaling - IN my resources when needed.

**Option A is incorrect.** Although going global is an advantage of using cloud due to availability of multiple regions, it does not account for our scenario

**Option B is incorrect.** Speed and agility of a website depends on various factors like application architecture / design, application communication etc..Although a cloud can provide a lot of hosting methodologies like Static web hosting using S3, Using a CDN like CloudFront for performance improvement, it does not account for in our scenario

**Option C is incorrect.** This is not the best advantage that can be achieved in this scenario. This question does not mention anything about DNS names.

**Option D is CORRECT.** Using the cloud will have a direct impact on my costs of purchasing new servers each year to support the dynamic increase in load.

### Reference:

<https://www.testpretraining.com/tutorial/aws-cloud-practitioner/define-the-aws-cloud-and-its-value-proposition/>

Ask our Experts

 View Queries

Did you like this Question?



### Question 36

Unattempted

Domain: Cloud Concepts

Which of the following is the customer responsibility with respect to AWS shared responsibility model for cloud Sustainability?

- A. Delivering efficient servers
- B. Optimizing Workload & resource utilization      right
- C. Sourcing Renewable power
- D. Achieve high server utilization rates

### Explanation:

Correct Answer: B

With respect to AWS shared responsibility model for cloud Sustainability, AWS is responsible for the sustainability of the cloud while customers are responsible for sustainability in the cloud.

Following are the responsibilities,

#### 1. AWS Responsibility:

- 1. Delivering efficient servers
- 2. Shared infrastructure
- 3. Water stewardship
- 4. Sourcing renewable power.

#### 2. Customer Responsibility:

- 1. Optimizing workloads and resource utilization
- 2. Minimizing the total resources required to be deployed for workloads.

**Option A is incorrect** as Delivering efficient servers is AWS's responsibility & not customer responsibility.

**Option C is incorrect** as Sourcing Renewable power is AWS's responsibility & not customer responsibility.

**Option D is incorrect** as Achieving high server utilization rates is AWS's responsibility & not customer responsibility.

For more information on Sustainability Pillar, refer to the following URL,

<https://docs.aws.amazon.com/wellarchitected/latest/sustainability-pillar/the-shared-responsibility-model.html>

<https://docs.aws.amazon.com/wellarchitected/latest/sustainability-pillar/design-principles-for-sustainability-in-the-cloud.html>

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?



### Question 37

Unattempted

Domain: Cloud Concepts

In the AWS Well-Architected Framework, which of the following is NOT a Security design principle to design solutions in AWS?

- A. Apply Security only at the edge of the network. right
- B. Protect Data at rest & in transit.
- C. Implement a strong Identity foundation.
- D. Enable Traceability.

### Explanation:

#### Correct Answer: A

Security needs to be applied at all network layers, like edge of network, VPC, all instances & application with the VPC. Applying Security controls at the edge of the network is not an efficient security control & against security design principles.

As per AWS Well-Architected Framework, the following are the design principles for security in the cloud:

- Implement a strong identity foundation.
- Enable traceability.

- Apply security at all layers.
- Automate security best practices.
- Protect data in transit and at rest.
- Keep people away from data.
- Prepare for security events.

**Options B, C, & D are incorrect** as these are part of security design principles that need to be followed while implementing security controls in the cloud.

For more information on Security Design Principle with AWS Well-Architected Framework, refer to the following URL:

<https://docs.aws.amazon.com/wellarchitected/latest/framework/sec-design.html>

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?  

### Question 38

Unattempted

Domain: Cloud Concepts

What is the key difference between an availability zone and an edge location?

- A. An availability zone is a grouping of AWS resources in a specific region. An edge location is a specific resource within the AWS region.
- B. An availability zone is an isolated location within an AWS region, whereas an edge location will deliver cached content to the closest location to reduce latency. right
- C. Edge locations are used as control stations for AWS resources.
- D. None of the above

### Explanation:

Answer – B

In AWS, there are regions with each region separated in a separate geographic area. Each region has multiple, isolated locations known as Availability Zones. An availability zone is used to host resources in a specific region.

For more information on Regions and availability zone, please visit the URL:

[Ask our Experts](#)[View Queries](#)

Did you like this Question?



### Question 39

Unattempted

**Domain:** Cloud Technology and Services

A financial company wants to automate the detection of suspicious banking transactions to alert customers. Which AWS service will meet this requirement?

- A. Amazon Rekognition
- B. Amazon Comprehend
- C. Amazon SageMaker      right
- D. Amazon Polly

### Explanation:

**Correct Answer: C**

**Option A is incorrect.** Amazon Rekognition is an ML service that helps perform image and video analysis within your applications. Examples are distinguishing real users from fake ones using facial verification in images and determining the similarity of a face image as compared with another image.

**Option B is incorrect.** Amazon Comprehend is a service that uses natural language processing (NLP) to find insights and relationships in text. Examples are Processing Financial Documents to classify & extract entities like “Insurance Claims”.

**Option C is CORRECT.** Amazon SageMaker is a fully managed AWS service that provides developers, and data scientists with the ability to build, train, and deploy machine learning (ML) models quickly. Once trained, it can be used for different solutions including “Detection of suspicious banking transactions” and alerting customers faster to avoid potential financial loss

**Option D is incorrect.** Amazon Polly is a service that turns text into speech and builds entirely new categories of speech-enabled products. As an example, your fitness band can tell you that it's now time to re-hydrate.

### References:

Use the following link to read more on different ML services offered by AWS including SageMaker

[Ask our Experts](#)[View Queries](#)

Did you like this Question?



## Question 40

Unattempted

**Domain:** Cloud Technology and Services

Which of the following statements related to CodePipeline are correct? (SELECT TWO)

- A. Only AWS-provided source repositories like CodeCommit can be used as the Source Provider in CodePipeline
- B. AWS CodePipeline can periodically check the selected code repository for code changes right
- C. One can use Jenkins as a Build Provider while configuring CodePipeline at the build stage right
- D. Both Build & Deploy stages are mandatory while configuring CodePipeline
- E. If a CodePipeline “Build” stage fails due to missing artifacts, one needs to trigger the pipeline manually after adding the artifacts to the code repository

## Explanation:

**Correct Answers:** B and C

**Option A is incorrect.** I can use different source providers like ‘Git repositories apart from CodeCommit as the code repository.

**Option B is CORRECT.** CodePipeline is automatically triggered when there is a change in the source repository e.g. I could have edited a source file. CodePipeline does this by periodically checking the selected code repository for any changes & triggers the pipeline once a change is detected.

**Option C is CORRECT.** Other than CodeBuild, I can configure third-party build tools like Jenkins for executing the Build stage of the pipeline

**Option D is incorrect.** A CodePipeline that is created should have at least two stages. The first stage is the “Source Stage”. The second one can be either a Build or a Deploy stage

**Option E is incorrect.** If a build fails due to missing artifacts, one can add them to the code repository & commit it. Once committed, CodePipeline will be triggered automatically beginning from the “Source Stage”

## References:

The reference mentioned below provides exhaustive documentation of CodePipeline

[What is AWS CodePipeline? - AWS CodePipeline \(amazon.com\)](#)

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?



### Question 41

Unattempted

**Domain:** Cloud Technology and Services

You have a Web application hosted in an EC2 Instance that needs to send notifications based on events. Which of the below services can assist in sending notifications?

- A. AWS SES
- B. AWS SNS right
- C. AWS SQS
- D. AWS EC2

### Explanation:

Answer – B

The AWS Documentation mentions the following:

Amazon Simple Notification Service (Amazon SNS) is a web service that enables applications, end-users, and devices to send and receive notifications from [the cloud](#) instantly.

For more information on AWS SNS, please visit the URL:

<https://aws.amazon.com/documentation/sns/>

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?



### Question 42

Unattempted

**Domain:** Cloud Concepts

Which AWS Well-Architected pillar focuses on using minimum required hardware for resources deployed in AWS Cloud ?

- A. Reliability Pillar
- B. Performance Efficiency Pillar
- C. Sustainability Pillar right
- D. Operational Excellence Pillar

### Explanation:

#### Correct Answer: C

Sustainability Pillar focuses on minimizing environment impact for running resources in AWS cloud. With respect to hardware needs , one of the best practices is to use minimum optimum amount of hardware for running resources in AWS cloud.

**Option A is incorrect** as the Reliability Pillar focuses on how quickly a resource recovers from a failure.

**Option B is incorrect** as the Performance Efficiency Pillar focuses on selecting optimised resource size for best performance.

**Option D is incorrect** as the Operational Excellence Pillar focuses on running & monitoring resources to meet business needs. It includes automating changes & defining SOP's (Standard Operating Procedures) for daily operations.

For more information on AWS Well-Architected and the Six Pillars, refer to the following URL,

<https://aws.amazon.com/architecture/well-architected/>

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?



### Question 43

Unattempted

**Domain:** Cloud Technology and Services

Which of the following services can be used to automate application code deployment to Amazon EC2 instance?

- A. AWS Elastic Beanstalk
- B. AWS CodeDeploy right
- C. AWS Config

## D. AWS CloudFormation

### Explanation:

**Correct Answer: B**

AWS Code Deploy can be used to automate deployment to Amazon EC2 instances and on-premise servers.

**Option A is incorrect** as AWS Elastic Beanstalk is an application management platform that can be used for managing application deployment handling capacity provisioning, load-balancing, auto-scaling & application health monitoring.

**Option C is incorrect** as AWS Config can be used to audit, evaluate configurations of AWS resources. It is not used to automate code deployment to Amazon EC2 instance.

**Option D is incorrect** as AWS CloudFormation is used for infrastructure deployment & automating resource creation within AWS Cloud Infrastructure.

For more information on AWS CodeDeploy, refer to the following URL:

<https://aws.amazon.com/codedeploy/faqs/?nc=sn&loc=6>

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?  

### Question 44

Unattempted

**Domain:** Cloud Technology and Services

Which of the following are benefits of the AWS's Relational Database Service (RDS)? Choose the 2 correct answers from the options below.

- A. Automated patches and backups right
- B. DB owner can resize the capacity accordingly right
- C. It allows you to store unstructured data
- D. It allows you to store NoSQL data

### Explanation:

Answer - A and B

The AWS Documentation mentions the following:

Amazon Relational Database Service (Amazon RDS) makes it easy to set up, operate, and scale a **relational database** in the cloud. It provides cost-efficient and resizable capacity while automating time-consuming administration tasks such as hardware provisioning, database setup, patching and backups. It frees you to focus on your applications. So you can give them the fast performance, high availability, security and compatibility they need.

For more information on AWS RDS, please visit the URL:

<https://aws.amazon.com/rds/>

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?



## Question 45

Unattempted

Domain: Security and Compliance

While using IAM within AWS, they recommend a list of best practices to be used while managing Users & their permissions. Which of the IAM practices are best to be chosen?

- A. An IAM user should be given default access to all services for being able to develop & deliver applications quickly
- B. Users should be provided both Passwords & Access keys within the AWS environment
- C. Within AWS, services should use roles for accessing other services right
- D. For cross account access it is best to share the Access Keys when one Account needs to access services of another Account

## Explanation:

Correct Answers: C

**Option A is incorrect.** The principle of “Grant Least Privilege” suggests providing the least & granular access to any User for accessing AWS services. This will help enhance the security of the system and prevent users from performing malicious/accidental tasks within the AWS environment that may compromise the system

**Option B is incorrect.** It's always best to see what kind of access the user requires. If a user does not use the Command Line Interface (CLI), it's best to avoid giving him access to Access Keys. Similarly, if a user only uses the CLI for development, he should be prohibited from accessing the AWS admin console using a password

**Option C is CORRECT.** Roles provide temporary access to AWS services & they don't require sharing the Access Keys. A role that is created may be assigned a definite set of permissions for

accessing different services. The service that uses the role actually assumes the role with which it can perform actions on another service. This way, we can induce granular permissions as well as prevent long-lived Access Keys from being compromised.

**Option D is incorrect.** Cross account access are scenarios where users in one Account would like to access services of other Accounts. Eg An user may use AWS Partner solutions to interact with a Customer's Account for monitoring it or do some orchestration. Access keys are long-lived credentials for which the AWS Partner must ensure to manage those keys by rotating them frequently to avoid them from being compromised. Roles defined with temporary access permissions when given to the user for accessing the Customer's account will be a best practice here without having to provide Access keys.

#### Reference:

<https://docs.aws.amazon.com/IAM/latest/UserGuide/best-practices.html>

[Ask our Experts](#)

 View Queries

Did you like this Question?



#### Question 46

Unattempted

**Domain:** Cloud Technology and Services

A company wants to create standard templates for the deployment of their Infrastructure. Which AWS service can be used in this regard?

- A. Amazon Simple Workflow Service
- B. AWS Elastic Beanstalk
- C. AWS CloudFormation right
- D. AWS OpsWorks

#### Explanation:

Answer – C

AWS CloudFormation gives developers and systems administrators an easy way to create and manage a collection of related AWS resources, provisioning and updating them in an orderly and predictable fashion.

For more information on Cloudformation, please visit the URL:

<https://aws.amazon.com/cloudformation/>

[Ask our Experts](#)[View Queries](#)

Did you like this Question?



### Question 47

Unattempted

**Domain:** Billing, Pricing, and Support

You have a distributed application that periodically processes large volumes of data across multiple Amazon EC2 Instances. The application is designed to recover gracefully from Amazon EC2 instance failures. You are required to accomplish this task most cost-effectively.

Which of the following will meet your requirements?

- A. Spot Instances right
- B. Reserved instances
- C. Dedicated instances
- D. On-Demand instances

### Explanation:

Answer - A

When you think of cost effectiveness, you can either have to choose Spot or Reserved instances. Now when you have a regular processing job, the best is to use spot instances. Since your application is designed to recover gracefully from Amazon EC2 instance failures, there is no issue even if you lose the Spot instances.

For more information on spot instances, please visit the URL:

<https://aws.amazon.com/ec2/spot/>

[Ask our Experts](#)[View Queries](#)

Did you like this Question?



Unattempted

**Question 48****Domain:** Cloud Technology and Services

You want to create a stream processing solution to process and query real-time streaming data using a SQL-based solution. You are looking for the simplest approach available that AWS provides.

What AWS service should you use?

- A. Amazon Kinesis Data Streams
- B. Amazon Kinesis Data Analytics right
- C. Amazon Kinesis Client Library
- D. Amazon Kinesis Data Firehose

**Explanation:****Answer: B**

**Option A is INCORRECT** because Amazon Kinesis Data Streams is not an AWS service to build analytics applications. It is an ingestion service that provides data streams to consumers (e.g., analytic applications).

AWS documentation mentions that “Amazon Kinesis Data Streams (KDS) is a massively scalable and durable real-time data streaming service.... The data collected is available in milliseconds to enable real-time analytics use cases such as real-time dashboards, real-time anomaly detection, dynamic pricing, and more”.

**Option B is CORRECT** because AWS documentation (FAQ section) mentions that “If you want a fully managed solution and you want to use SQL to process the data from your data stream, you should use Kinesis Data Analytics”.

**Option C is INCORRECT** because Amazon Kinesis Client Library (KCL) is a service to process and query streaming data using SQL. KCL is more complicated than AWS Kinesis Data Analytics.

AWS documentation mentions that “Kinesis Data Analytics uses the KCL to read data from streaming data sources as one part of your underlying application. The service abstracts this from you, as well as many of the more complex concepts associated with using the KCL, such as checkpointing”.

**Option D is INCORRECT** because Amazon Kinesis Data Firehose is a service for loading data streams and not creating SQL-based applications.

**Diagram:** none**References:**

<https://aws.amazon.com/kinesis/data-streams/>

<https://aws.amazon.com/kinesis/data-analytics/faqs/?nc=sn&loc=6> (explanation of option C)

<https://aws.amazon.com/kinesis/data-firehose/>

Ask our Experts

 View Queries

Did you like this Question?



### Question 49

Unattempted

Domain: Cloud Technology and Services

What of the following AWS services allows developers to deploy and manage web applications on the cloud easily?

- A. Amazon Route 53
- B. AWS Elastic Beanstalk right
- C. Amazon CloudFront
- D. Elastic Container Service

### Explanation:

#### Answer - B

AWS Elastic Beanstalk makes it even easier for developers to deploy and manage web applications in the AWS Cloud quickly. Developers simply upload their application and Elastic Beanstalk automatically handles the deployment details of capacity provisioning, load balancing, auto-scaling, and application health monitoring.

For more information on Elastic Beanstalk, please visit the URL:

<https://aws.amazon.com/elasticbeanstalk/faqs/>

Ask our Experts

 View Queries

Did you like this Question?



### Question 50

Unattempted

Domain: Cloud Technology and Services

A company is deploying a new two-tier web application in AWS. The company wants to store their most frequently used data to improve the response time for the application. Which AWS service provides the caching solution for the company's requirements?

A. MySQL Installed on two Amazon EC2 Instances in a single Availability Zone

B. Amazon RDS for MySQL with Multi-AZ

C. Amazon ElastiCache right

D. Amazon DynamoDB

### Explanation:

Answer – C

Amazon ElastiCache is a web service that makes it easy to deploy, operate, and scale an in-memory data store or cache in the cloud. The service improves the performance of web applications by allowing you to retrieve information from fast, managed, in-memory data stores, instead of relying entirely on slower disk-based databases.

For more information on Elastic cache, please visit the link:

<https://aws.amazon.com/elasticache/>

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?  

### Question 51

Unattempted

Domain: Cloud Concepts

If you want to take a backup of an EBS Volume, what would you do?

A. Store the EBS volume in S3.

B. Store the EBS volume in an RDS database.

C. Create an EBS snapshot. right

D. Store the EBS volume in DynamoDB.

### Explanation:

Answer - C

The AWS Documentation mentions the following:

You can back up the data on your Amazon EBS volumes to Amazon S3 by taking point-in-time snapshots.

For more information on EBS Snapshots, please visit the link:

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>

[Ask our Experts](#)

 View Queries

Did you like this Question?



## Question 52

Unattempted

**Domain:** Cloud Technology and Services

Which of the following components of the AWS Batch can be used to define the memory requirement of the resources ?

- A. Jobs
- B. Compute Environment
- C. Job Queue
- D. Job Definitions      right

### Explanation:

**Correct Answer: D**

Job Definitions can be used to specify how a job runs in a compute environment. It can be used to define memory, CPU requirements , IAM roles required for running the Jobs, container properties, environment variables, and mount points for persistent storage

**Option A is incorrect** as Jobs is a script or executable file which is submitted to AWS Batch. Parameters for running a Job are defined in Job Definitions.

**Option B is incorrect** as Compute Environment is a set of compute resources which are used to run the Jobs. These resources can be either managed compute resources or unmanaged compute resources.

**Option C is incorrect** as Job Queue can be used to specify priority of the jobs which are submitted to AWS Batch. For time-sensitive jobs , high priority can be defined while low priority can be assigned to jobs which can run at off-peak hours when resources are cheaper.

For more information on components of AWS Batch, refer to the following URL,

[Ask our Experts](#)[View Queries](#)

Did you like this Question?



### Question 53

Unattempted

**Domain:** Cloud Concepts

Which of the following options of AWS RDS allows for AWS to failover to a secondary database in case the primary one fails?

- A. Amazon RDS Multi-AZ deployments right
- B. AWS Failover
- C. AWS Secondary
- D. AWS Standby

### Explanation:

#### Answer – A

The AWS Documentation mentions the following.

Amazon RDS Multi-AZ deployments provide enhanced availability and durability for Database (DB) Instances, making them a natural fit for production database workloads. When you provision a Multi-AZ DB Instance, Amazon RDS automatically creates a primary DB Instance and synchronously replicates it to a standby instance in a different Availability Zone (AZ). Each AZ runs on its own physically distinct, independent infrastructure and is engineered to be highly reliable. In case of an infrastructure failure, Amazon RDS performs an automatic failover to the standby (or to a read replica in the case of Amazon Aurora). So you can resume database operations as soon as the failover is complete.

**Options B, C and D** are the distractors because there is no AWS service like failover, standby or secondary. Failover is the situation when Multi-AZ appears and Standby or Secondary represents other RDS servers appears during Multi-AZ deployments.

For more information on AWS RDS Multi-AZ, visit the below link:

<https://aws.amazon.com/rds/details/multi-az/>

[Ask our Experts](#)[View Queries](#)

Did you like this Question?



### Question 54

Unattempted

**Domain:** Cloud Concepts

For which of the following AWS resources, the Customer is responsible for the infrastructure-related security configurations?

- A. Amazon RDS
- B. Amazon DynamoDB
- C. Amazon EC2      right
- D. AWS Fargate

### Explanation:

**Correct Answer: C**

Amazon EC2 is an Infrastructure as a Service (IaaS) for which customers are responsible for the security and the management of guest operating systems.

**Options A, B, and D are incorrect** as all these resources are part of abstracted services for which AWS is responsible for the security, & infrastructure layer. Customers are responsible for data that is saved on these resources.

For more information on the Shared responsibility model, refer to the following URL:

<https://aws.amazon.com/compliance/shared-responsibility-model/>

[Ask our Experts](#)[View Queries](#)

Did you like this Question?



### Question 55

Unattempted

**Domain:** Cloud Technology and Services

Which of the following statements regarding Amazon Athena are accurate? (Select TWO)

- A. Amazon Athena queries data directly from Amazon S3.      right

- B. Amazon Athena is not suitable for complex analysis such as large joins, window functions and arrays.
- C. Amazon Athena resources are allocated in accordance to processing and memory requirements prior to deployment.
- D. Amazon Athena is compatible with data formats such as CSV, JSON, ORC, AVRO and Parquet right
- E. Amazon Athena uses a variety of query languages including SQL, LDAP, JPQL as well as CQL.

### Explanation:

Correct Answer – A, D

Amazon Athena is a serverless query service that does not need to build databases on dedicated Elastic Block Store (EBS) volumes. Instead, it builds tables from data read directly from Amazon S3 buckets. Amazon Athena does not store any of the data. The service is compatible with the regular data formats that include CSV, JSON, ORC, AVRO and Parquet.

<https://docs.aws.amazon.com/athena/latest/ug/what-is.html>

**Option B is incorrect** because Amazon Athena can query Big Data, complex analysis such as large joins, window functions and arrays.

**Option C is incorrect** because Amazon Athena is serverless. Thus the service scales following the resource demands. No prior resource planning is necessary.

**Option E is incorrect** because Amazon Athena uses SQL only.

Ask our Experts

 View Queries

Did you like this Question?



### Question 56

Unattempted

Domain: Cloud Technology and Services

An administrator is running a large deployment of AWS resources that are spread across several AWS Regions. They would like to keep track of configuration changes on all the resources and maintain a configuration inventory. What is the best service they can use?

- A. AWS CloudFormation
- B. Stacks and Templates

C. AWS Backup

D. AWS Config right

### Explanation:

#### Correct Answer – D

AWS Config will meet the scenario requirements. The service allows the administrator to monitor and record configuration changes on AWS resources in their account. The service also allows the administrator to create a resource configuration inventory.

<https://aws.amazon.com/config/>

**Option A is incorrect** because AWS CloudFormation will allow the administrator to create templates of resources such as EC2 instances and RDS instances but not the actual configurations in these resources.

**Option B is incorrect** because Templates and Stacks form the basis of AWS CloudFormation. They aid in the automated deployment of whole environments but not the applications that run in them.

**Option C is incorrect** because AWS Backup is a fully managed service that allows the administrator to back up data in the cloud and on-premises. The service is not the most appropriate to monitor and record resource configuration changes.

[Ask our Experts](#)

 View Queries

Did you like this Question?



### Question 57

Unattempted

**Domain:** Cloud Technology and Services

Whilst working on a collaborative project, an administrator would like to record the initial configuration and several authorized changes that engineers make to the route table of a VPC. What is the best method to achieve this?

A. Use of AWS Config right

B. Use of VPC Flow Logs

C. Use of AWS CloudTrail

D. Use of an AWS Lambda function that is triggered to save a log file to an S3 bucket each time configuration changes are made.

## Explanation:

### Correct Answer – A

AWS Config can be used to keep track of configuration changes on AWS resources, keeping multiple date-stamped versions in a reviewable history. This makes it the best method to meet the scenario requirements.

<https://aws.amazon.com/config/>

**Option B is incorrect** because VPC flow logs will only capture IP traffic-related information passing through and from network interfaces within the VPC. VPC flow logs will not be able to capture configuration changes made to route tables.

<https://docs.aws.amazon.com/vpc/latest/userguide/flow-logs.html>

**Option C is incorrect** because AWS CloudTrail will capture identity access activity, event history into the AWS environment. Recording the actions and API calls are not best suited to keep a record of configurations.

<https://aws.amazon.com/cloudtrail/>

**Option D is incorrect** because using a Lambda function to write configuration changes might meet the requirements, but it would not be the best method. AWS Config can deliver what is needed with much less administrative input.

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?



## Question 58

Unattempted

**Domain:** Cloud Technology and Services

A company does not want to manage their database. Which of the following services is a fully managed NoSQL database provided by AWS?

- A. AWS RDS
- B. DynamoDB right
- C. Oracle RDS
- D. Elastic Map Reduce

## Explanation:

Answer – B

DynamoDB is a fully managed NoSQL offering provided by AWS. It is now available in most regions for users to consume.

For more information on AWS DynamoDB, visit the below link:

<https://aws.amazon.com/dynamodb/>

[Ask our Experts](#)

 [View Queries](#)

Did you like this Question?



### Question 59

Unattempted

**Domain:** Billing, Pricing, and Support

Which of the following tools can be used to create an estimated cost for a new solution to be deployed on AWS Cloud infrastructure?

- A. AWS Cost and Usage Report
- B. AWS Budgets
- C. AWS Cost Explorer
- D. AWS Pricing Calculator right

### Explanation:

**Correct Answer : D**

AWS Pricing Calculator helps create estimated usage costs for all resources to be deployed on AWS Cloud Infrastructure.

**Option A is incorrect** as AWS Cost and Usage Report will provide detailed data for usage summary.

**Option B is incorrect** as AWS Budget can set a custom budget to track usage for AWS resources.

**Option C is incorrect** as AWS Cost Explorer can be used to analyze & manage AWS usage.

For more information on AWS Pricing Calculator, refer to the following URL:

<https://docs.aws.amazon.com/pricing-calculator/latest/userguide/what-is-pricing-calculator.html>

[Ask our Experts](#)[View Queries](#)

Did you like this Question?



## Question 60

Unattempted

**Domain:** Cloud Technology and Services

Which of the following services helps provide a connection from on-premises infrastructure to resources hosted in the AWS Cloud? Choose 2 answers from the options given below.

- A. AWS VPC
- B. AWS VPN      right
- C. AWS Direct Connect      right
- D. AWS Subnets

### Explanation:

Answer – B and C

The AWS Documentation mentions the following.

AWS Direct Connect makes it easy to establish a dedicated network connection from your premises to AWS. Using AWS Direct Connect, you can establish private connectivity between AWS and your datacenter, office, or colocation environment, which in many cases can reduce your network costs, increase bandwidth throughput, and provide a more consistent network experience than Internet-based connections.

You can connect your VPC to remote networks by using a VPN connection.

For more information on AWS Direct Connect, please visit the link:

<https://aws.amazon.com/directconnect/?p=tile>

For more information on AWS VPN, please visit the link:

<https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpn-connections.html>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?

**Question 61**

Unattempted

**Domain:** Cloud Technology and Services

A small finance institute is looking for an option of launching an accounting application in AWS cloud. Which of the following services can be used for quick launch with a pre-configured resource?

- A. Amazon Workspaces
- B. AWS Elastic Beanstalk
- C. AWS Batch
- D. Amazon LightSail      right

**Explanation:****Correct Answer: D**

Amazon LightSail is a virtual private server which can be used to quickly launch business applications in AWS cloud at a low cost .It can also be used to launch simple web sites and to set up a test environment. These applications are launched on a pre-configured resource in an AWS cloud.

**Option A is incorrect** as Amazon Workspaces provides a secure managed service for virtual desktops for remote users. It is not suitable for launching a quick business application..

**Option B is incorrect** as AWS Elastic Beanstalk deploys web applications on AWS infrastructure. Once application code is uploaded, AWS Elastic Beanstalk handles deployment along with capacity provisioning, auto-scaling, load balancing & health monitoring. Amazon LightSail is better suited for simple applications and when flexibility in resource configurations is not required.

**Option C is incorrect** as AWS Batch can be used for managed batch processing activities using compute like Amazon EC2,AWS Fargate or Spot instance.

For more information on Amazon LightSail, refer to the following URL,

<https://aws.amazon.com/lightsail/>

[Ask our Experts](#)[View Queries](#)

Did you like this Question?



**Question 62**

Unattempted

**Domain:** Security and Compliance

Which AWS Service provides a centralized platform for aggregating, prioritizing, and managing security alerts and compliance findings across an organization's AWS environment?

- A. AWS Shield
- B. Amazon GuardDuty
- C. AWS Security Hub right
- D. AWS WAF

**Explanation:****Correct Answer: C**

AWS Security Hub is a central security tool to manage security across several AWS accounts and automate security checks. Security Hub enables you to have a comprehensive view of your security posture and compliance status, making it easier to identify and address potential security issues and compliance violations.

**Option A is incorrect** because AWS Shield is a managed Distributed Denial of Service (DDoS) protection service provided by AWS. It helps safeguard applications running on AWS against DDoS attacks by automatically detecting and mitigating malicious traffic.

**Option B is incorrect** because Amazon GuardDuty is a threat detection service that helps protect AWS accounts and workloads by continuously monitoring for malicious activity and unauthorized behavior. It generates findings and alerts when any suspicious activity is detected.

**Option D is incorrect** because AWS WAF is a managed security service that helps protect web applications from various types of online threats and attacks by allowing you to define and enforce custom rules for your web traffic.

**Reference:**

<https://docs.aws.amazon.com/securityhub/latest/userguide/what-is-securityhub.html>

**Ask our Experts****(+)** View Queries

Did you like this Question?

**Question 63**

Unattempted

**Domain:** Billing, Pricing, and Support

There is a requirement to host a set of servers in the Cloud for a short period of 3 months. The instances should be highly available and cost-effective. The instances can not be interrupted during processing. Which of the following types of instances should be chosen?

- A. Spot Instances
- B. On-Demand right
- C. No Upfront costs Reserved
- D. Partial Upfront costs Reserved

### Explanation:

Answer - B

Since the requirement is just for 3 months and the instances cannot be interrupted during processing, the most suitable type is On-Demand Instances.

For more information on AWS On-Demand Instances, please visit the link:

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-on-demand-instances.html>

[Ask our Experts](#)

 View Queries

Did you like this Question?



### Question 64

Unattempted

Domain: Security and Compliance

Which of the below can be configured to enhance the security at the subnet level?

- A. Virtual Private Cloud (VPC)
- B. Configure transitive VPC peering
- C. NACL (Network Access Control List) right
- D. Security Group

### Explanation:

Answer: C

**Option A is INCORRECT.** Virtual Private Cloud (VPC) is a virtual network that lets us launch AWS resources in the defined virtual network.

**Option B is INCORRECT.** Configure transitive VPC peering is invalid as this is not supported in AWS.

**Option C is CORRECT.** NACLs can be configured to enhance the security at the subnet level.

**Option D is INCORRECT.** Security Group acts as a virtual firewall by controlling the traffic both inbound and outbound. Security group acts at the instance level.

#### Reference:

<https://docs.aws.amazon.com/vpc/latest/userguide/what-is-amazon-vpc.html>

<https://docs.aws.amazon.com/vpc/latest/peering/invalid-peering-configurations.html>

<https://docs.aws.amazon.com/vpc/latest/userguide/vpc-network-acls.html>

[https://docs.aws.amazon.com/vpc/latest/userguide/VPC\\_SecurityGroups.html](https://docs.aws.amazon.com/vpc/latest/userguide/VPC_SecurityGroups.html)

[Ask our Experts](#)

 [View Queries](#)

Did you like this **Question?**



#### Question 65

Unattempted

**Domain:** Cloud Concepts

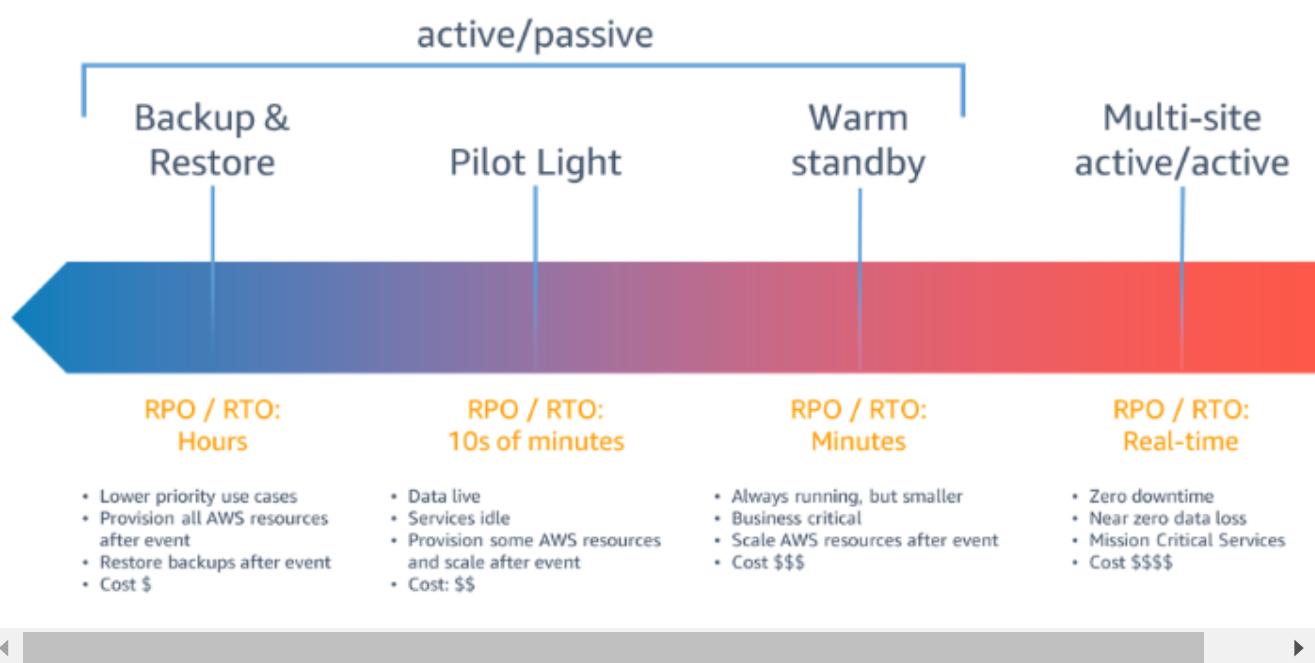
Which of the following is not a disaster recovery deployment technique?

- A. Pilot light
- B. Warm standby
- C. Single Site right
- D. Multi-Site

#### Explanation:

Answer - C

The below snapshot from the AWS documentation shows the different disaster recovery techniques.



For more information on Disaster recovery techniques, please refer to the below URL:

<https://docs.aws.amazon.com/whitepapers/latest/disaster-recovery-workloads-on-aws/disaster-recovery-options-in-the-cloud.html>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?



[Finish Review](#)



[Hands-on Labs](#)   [Sandbox](#)   [Pricing](#)   [For Business](#)   [Library](#)

## Categories

- Cloud Computing Certificatio...
- Amazon Web Services (AWS)
- Microsoft Azure
- Google Cloud
- DevOps
- Cyber Security
- Microsoft Power Platform
- Microsoft 365 Certifications
- Java Certifications

## Popular Courses

- AWS Certified Solutions Archit...
- AWS Certified Cloud Practition...
- Microsoft Azure Exam AZ-204 ...
- Microsoft Azure Exam AZ-900 ...
- Google Cloud Certified Associ...
- Microsoft Power Platform Fund...
- HashiCorp Certified Terraform...
- Snowflake SnowPro Core Certi...
- Docker Certified Associate

## Company

- About Us
- Blog
- Reviews
- Careers
- Become an Affiliate
- Become Our Instructor
- Team Account
- AWS Consulting Services

**Legal**

Privacy Policy  
Terms of Use  
EULA  
Refund Policy  
Programs Guarantee

**Support**

Contact Us  
Discussions  
FAQs

Need help? Please  or  +91 6364678444



---

©2023, Whizlabs Software Pvt. Ltd. All rights reserved.